

**AIM: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.****Create 3 EC2 Ubuntu Instances on AWS.**

Login to your AWS console. Go to services and in that search for EC2 and create 3 EC2 Ubuntu Instances as master 1, node1 and node 2. While making an instance make sure to select Amazon Linux and in linux type instead of default t2.micro select t2.medium.

**Name and tags** [Info](#)

Name

master1

Add additional tags

▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Q Search our full catalog including 1000s of application and OS images

Recents Quick Start

Amazon Linux

macOS

Ubuntu

Windows

Red Hat

SUSE Li

Browse more AMIs  
Including AMIs from AWS, Marketplace and

▼ Instance type [Info](#) | [Get advice](#)

Instance type

t2.medium

Family: t2 2 vCPU 4 GiB Memory Current generation: true  
On-Demand Linux base pricing: 0.0464 USD per Hour  
On-Demand RHEL base pricing: 0.0752 USD per Hour  
On-Demand Windows base pricing: 0.0644 USD per Hour  
On-Demand SUSE base pricing: 0.1464 USD per Hour

All generations

[Compare instance types](#)

Additional costs apply for AMIs with pre-installed software

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

vockey

Create new key pair

## Setting SSH for establishing connections

[illegible]

## INSTALLATION OF DOCKER

For installing docker we use the following steps:

**STEP 1:**In node 1 EC2 instance install docker and repeat the same step for master and node2 .

Syntax: `yum install docker -y`

```

root@ip-172-31-176 ec2-user]# yum install docker -y
Last metadata expiration check: 0:21:25 ago on Fri Sep 13 17:05:55 2024.
Dependencies resolved.

```

Package	Architecture	Version	Repository	Size
<b>Installing:</b>				
docker	x86_64	25.0.6-1.amzn2023.0.2	amazonlinux	44 M
<b>Installing dependencies:</b>				
containerd	x86_64	1.7.20-1.amzn2023.0.1	amazonlinux	35 M
iptables-libs	x86_64	1.8.8-3.amzn2023.0.2	amazonlinux	401 k
iptables-nft	x86_64	1.8.8-3.amzn2023.0.2	amazonlinux	183 k
libcgroupp	x86_64	3.0-1.amzn2023.0.1	amazonlinux	75 k
libnetfilter_conntrack	x86_64	1.0.8-2.amzn2023.0.2	amazonlinux	58 k
libnftnl	x86_64	1.0.1-19.amzn2023.0.2	amazonlinux	30 k
libnftnl	x86_64	1.2.2-2.amzn2023.0.2	amazonlinux	84 k
pigs	x86_64	2.5-1.amzn2023.0.3	amazonlinux	83 k
runsc	x86_64	1.1.13-1.amzn2023.0.1	amazonlinux	3.2 M

```

Transaction Summary
Install 10 Packages

```

i-0defb5859fc2b0488 (node1)
PublicIPs: 54.157.60.252 PrivateIPs: 172.31.21.176

**STEP 2:** After the installation of docker is successfully completed in all the three instances start the docker by the syntax given below:

Syntax :systemctl start docker.

Start the docker in master and node2 too .

```

Running scriptlet: docker-25.0.6-1.amzn2023.0.2.x86_64
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.

Verifying      : containerd-1.7.20-1.amzn2023.0.1.x86_64          1/10
Verifying      : docker-25.0.6-1.amzn2023.0.2.x86_64           2/10
Verifying      : iptables-libs-1.8.8-3.amzn2023.0.2.x86_64      3/10
Verifying      : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64      4/10
Verifying      : libcgroupp-3.0-1.amzn2023.0.1.x86_64          5/10
Verifying      : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64 6/10
Verifying      : libnftnl-1.0.1-19.amzn2023.0.2.x86_64         7/10
Verifying      : libnftnl-1.2.2-2.amzn2023.0.2.x86_64          8/10
Verifying      : pigz-2.5-1.amzn2023.0.3.x86_64                9/10
Verifying      : runc-1.1.13-1.amzn2023.0.1.x86_64             10/10

Installed:
  containerd-1.7.20-1.amzn2023.0.1.x86_64    docker-25.0.6-1.amzn2023.0.2.x86_64    iptables-libs-1.8.8-3.amzn2023.0.2.x86_64
  iptables-nft-1.8.8-3.amzn2023.0.2.x86_64    libcgroupp-3.0-1.amzn2023.0.1.x86_64    libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64
  libnftnl-1.0.1-19.amzn2023.0.2.x86_64      libnftnl-1.2.2-2.amzn2023.0.2.x86_64    pigz-2.5-1.amzn2023.0.3.x86_64
  runc-1.1.13-1.amzn2023.0.1.x86_64

Complete!
[root@ip-172-31-21-176 ec2-user]# systemctl start docker
[root@ip-172-31-21-176 ec2-user]#

```

i-0defb5859fc2b0488 (node1)

PublicIPs: 54.157.60.252 PrivateIPs: 172.31.21.176

## INSTALLATION OF KUBERNETES

After installing and starting the docker in all the three instances ,now lets install kubernetes for the installation we use the following steps:

**STEP 1:**Set SELinux to permissive mode:

Syntax:sudo setenforce 0

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/'
/etc/selinux/config
```

```
[root@ip-172-31-25-172 docker]# sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

**STEP 2:**Add the Kubernetes yum repository. The exclude parameter in the repository definition ensures that the packages related to Kubernetes are not upgraded upon running yum update as there's a special procedure that must be followed for upgrading Kubernetes

```
[root@ip-172-31-21-176 ec2-user]# sudo su
[root@ip-172-31-21-176 ec2-user]# yum repolist
repo id                                repo name
amazonlinux                            Amazon Linux 2023 repository
kernel-livepatch                       Amazon Linux 2023 Kernel Livepatch repository
[root@ip-172-31-21-176 ec2-user]# cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[root@ip-172-31-21-176 ec2-user]#
```

i-0defb5859fc2b0488 (node1)

PublicIPs: 54.157.60.252 PrivateIPs: 172.31.21.176

### STEP 3: Install kubelet, kubeadm and kubectl:

Syntax: `sudo yum install -y kubelet kubeadm kubectl`  
`--disableexcludes=kubernetes`

```
Last login: Fri Sep 13 17:58:28 2024 from 18.206.107.27
[ec2-user@ip-172-31-21-176 ~]$ sudo su
[root@ip-172-31-21-176 ec2-user]# sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Kubernetes
Dependencies resolved.
60 kB/s | 9.4 kB    00:00
=====
Package                                Architecture      Version           Repository        Size
-----
Installing:
kubeadm                                x86_64            1.31.1-150500.1.1  kubernetes        11 M
kubectl                                x86_64            1.31.1-150500.1.1  kubernetes        11 M
kubelet                                x86_64            1.31.1-150500.1.1  kubernetes        15 M
Installing dependencies:
conntrack-tools                        x86_64            1.4.6-2.amzn2023.0.2  amazonlinux        208 k
cri-tools                              x86_64            1.31.1-150500.1.1  kubernetes         6.9 M
kubernetes-cni                         x86_64            1.5.1-150500.1.1  kubernetes         7.1 M
libnetfilter_cthelper                  x86_64            1.0.0-21.amzn2023.0.2  amazonlinux        24 k
libnetfilter_cttimeout                 x86_64            1.0.0-19.amzn2023.0.2  amazonlinux        24 k
libnetfilter_queue                     x86_64            1.0.5-2.amzn2023.0.2  amazonlinux        30 k
Transaction Summary
-----
Install 9 Packages
```

### STEP 4: Enable the kubelet service before running kubeadm:

Syntax: `sudo systemctl enable --now kubelet`

```
[root@ip-172-31-21-176 ec2-user]# sudo systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service.
[root@ip-172-31-21-176 ec2-user]#
```

i-0defb5859fc2b0488 (node1)

PublicIPs: 54.157.60.252 PrivateIPs: 172.31.21.176

**STEP 5:** It can be seen from the `repolist` command which lists all the repository we can see that `kubernetes` is installed repeat all these steps on `master1` and `node2`.

```
[root@ip-172-31-21-176 ec2-user]# yum repolist
repo id                                repo name
amazonlinux                            Amazon Linux 2023 repository
kernel-livepatch                       Amazon Linux 2023 Kernel Livepatch repository
kubernetes                             Kubernetes
[root@ip-172-31-21-176 ec2-user]#
```

i-0defb5859fc2b0488 (node1)

PublicIPs: 54.157.60.252 PrivateIPs: 172.31.21.176

**STEP 6 :** This command disable swap space and configure the system to use iptables for bridged network traffic, then apply these settings.

Syntax: `sudo swapoff -a`

`echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf`  
`sudo sysctl -p`

```
[root@ip-172-31-16-56 ec2-user]# sudo swapoff -a
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
sudo sysctl -p
net.bridge.bridge-nf-call-iptables=1
```

## STEP 7: Initialize Kubernetes in master instance .

Syntax: `kubeadm init`

```
[root@ip-172-31-16-56 ec2-user]# kubeadm init
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[WARNING FileExisting-socat]: socat not found in system path
[WARNING FileExisting-tc]: tc not found in system path
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W0913 18:58:26.902514 34809 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent
with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.10" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-16-56.ec2.internal kubernetes kubernetes.default kubernetes.default.svc kubernetes.default.svc.cluster.local] and IPs [10.96.0.1 172.31.16.56]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-16-56.ec2.internal localhost] and IPs [172.31.16.56 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-16-56.ec2.internal localhost] and IPs [172.31.16.56 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key

i-0ddf50a232db19957 (master1 )
PublicIPs: 3.88.204.138 PrivateIPs: 172.31.16.56
```

```
To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.16.56:6443 --token oghyi3.fnspro8pevgr0d5 \
--discovery-token-ca-cert-hash sha256:ec71ffc0d9fd79263fb8909d938da8d29e5f15a21ab5e0a17ec93514e8c4ecb8
```

## STEP 8: Use the mkdir and chown commands shown above

```
[root@ip-172-31-16-56 ec2-user]# mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

**STEP 9:**Add a common networking plugin called flannel

Syntax:kubectl apply -f

<https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

```
[root@ip-172-31-16-56 ~]# kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
```

**STEP 10:**Apply deployment of nginx server using the following command.

Syntax:

kubectl apply -f <https://k8s.io/examples/application/deployment.yaml>

```
[root@ip-172-31-16-56 ~]# kubectl apply -f https://k8s.io/examples/application/deployment.yaml
deployment.apps/nginx-deployment created
```

**STEP 11:** Next copy and past the join link in the worker nodes so that the worker nodes can join the cluster.

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.82.191:6443 --token 8450pt.tdprcovwa6lrgyol \
--discovery-token-ca-cert-hash sha256:b11f191f3df19a2e9112a5c19b4461bffeadd8b5be8625ad8451019aecc043c
```

**STEP 12:** We can check the nodes that have joined the cluster using kubectl get nodes. Right now there is only one node which is the master node.

```
[root@ip-172-31-85-89 ec2-user]# kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-85-89.ec2.internal	NotReady	control-plane	72s	v1.26.0

**STEP 13:** After performing join commands on the worker nodes we will get following output:

```
This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

Once again when you run kubectl get nodes you will now see all 3 nodes have joined the cluster:

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-85-89.ec2.internal	NotReady	control-plane	119s	v1.26.0
ip-172-31-89-46.ec2.internal	NotReady	<none>	19s	v1.26.0
ip-172-31-94-70.ec2.internal	NotReady	<none>	12s	v1.26.0

**Conclusion:** In this experiment we have created 3 EC2 instances, while making instance make sure to click on AmazonLinux and change the instance type to t3.medium or large if it says the memory space or number of CPU's is not enough. Setting SSH for establishing connections in that we have installed and started docker and kubernetes ,initialising kubernetes and by performing various steps we have learned how to link both of the nodes that is node1 and node2 to the main node that is master1 .